

Nicholas C. Rowley (SBN 220036)
Lead Counsel
TRIAL LAWYERS FOR JUSTICE
8235 Santa Monica Blvd.
West Hollywood, CA 90046
Tel.: (310) 855-3727
Fax: (310) 855-3595
ncr@tl4j.com

Andrew T. Tutt (SBN 343206)
R. Stanton Jones (DC Bar No. 987088)*
TRIAL LAWYERS FOR JUSTICE
421 W. Water Street
Decorah, IA 52101
Tel.: (202) 490-7080
Fax: (563) 204-4205 (Tutt)
andrew@tl4j.com; stanton@tl4j.com

Jakob Z. Norman (WY Bar No. 6-3685)*
TRIAL LAWYERS FOR JUSTICE
724 S. Durbin Street
Casper, WY 82601
Tel.: (307) 288-4400 ext. 101
jakob@tl4j.com

Attorneys for Plaintiffs and the Proposed Class

* Application for *Pro Hac Vice* Admission Forthcoming

**IN THE UNITED STATES DISTRICT COURT
FOR THE NORTHERN DISTRICT OF CALIFORNIA
SAN FRANCISCO DIVISION**

CHARLES BUIST, CHEYENNE HUNT,
CHRISTINE BULLOCK, and NICK
SPETSAS, individually and on behalf of all
others similarly situated,

Plaintiffs,

v.

ANTHROPIC, PBC; OPENAI OPCO, LLC;
SPACEXAI LLC; AND GOOGLE LLC,

Defendants.

Case No. 3:26-cv-10693

COMPLAINT

Class Action

Demand for Jury Trial



INTRODUCTION

1. This action challenges an agreement among the four companies at the frontier of artificial intelligence—Anthropic, OpenAI, SpaceXAI, and Google—to slow the pace at which each company improves the competing products it sells to consumers.

2. The agreement was proposed in public, accepted in public, and confirmed in public. On the morning of September 12, 2026, Anthropic’s Chief Executive Officer, Dario Amodei, published an essay titled *We Must Pace the Frontier*. It declared: “We must slow the pace at which we improve the capabilities of AI models.” It called for “industry-wide coordination” among frontier laboratories to set “limits on the rate of unchecked AI progress,” and it explained that coordinated slowing would let each participant reduce its pace “without sacrificing commercial advantage.”

3. Amodei’s rivals confirmed their agreement the same day. Within about an hour, Elon Musk, who founded and controls SpaceXAI’s Grok business, quote-posted the proposal and wrote: “Dario is right.” OpenAI’s Chief Executive Officer, Sam Altman, wrote: “I agree with Dario that we need to pace the frontier,” and committed OpenAI to the first concrete step of Amodei’s plan. Google DeepMind co-founder Demis Hassabis called the essay “the right path forward” and tied it to the “industry-wide standards body for frontier AI” he had proposed two months earlier as a vehicle through which the laboratories could coordinate.

4. Two days later, Altman said what pacing meant: AI progress “should be slower than it otherwise could be.” He added that OpenAI would not wait for an antitrust exemption or legislation before beginning that work with its “colleagues across the industry.” The following day, OpenAI’s Global Policy Chief confirmed that OpenAI, Anthropic, and Google DeepMind had already been working together for weeks, reportedly through a working group of company representatives that had met regularly since July.

5. Defendants themselves have explained why they needed an agreement. Frontier-AI companies compete above all by improving their models faster than their rivals: better reasoning, better coding, more capable agents, larger context windows, new features. A July 2026 statement signed by Amodei, OpenAI’s Chief Scientist, a Google DeepMind co-founder, and dozens of their



1 colleagues acknowledged that each company faces “intense competitive pressure not to unilaterally
2 ally slow” development. A firm that slows alone loses customers, revenue, and technological leadership
3 to rivals that keep going. Amodei’s proposal solved that problem the way cartels always
4 have: by agreeing to slow together, so that no participant bears the competitive cost of restraint.

5 6. An agreement among competitors to reduce the quality of their products and the rate at
6 which those products improve is an agreement to restrict output. It is among the restraints Section
7 1 of the Sherman Act, 15 U.S.C. § 1, has always condemned, and it does not become lawful because
8 the products are new.

9 7. Plaintiffs take AI safety seriously. Plaintiffs, and each of them, believe that AI should
10 have guardrails with ethical considerations and direct liability for any breach being paramount.
11 Those guardrails and consequences should not be decided by the Defendants, but by we the people,
12 through state and federal regulation, and through the voices of juries in our communities.

13 8. Artificial intelligence is among the most consequential technologies in human history.
14 Researchers have used it to predict the three-dimensional structure of nearly every protein known
15 to science, to identify cancers from imaging studies with accuracy matching or exceeding trained
16 specialists, to accelerate the development of new materials for clean energy, to solve previously
17 unsolved mathematical problems, and to compress years of scientific labor into days. It is undisputed
18 that AI alone could add trillions of dollars annually to the global economy. These are not
19 hypothetical gains. They are the product of the competitive race that Defendants’ agreement now
20 threatens to slow.

21 9. Plaintiffs recognize that powerful technology carries real risks, and they take those risks
22 seriously. But the lesson of every previous transformative industry is that safety and responsibility
23 are properly the work of each company acting on its own, not a cartel license. Automobile manufacturers
24 do not allow ordinary passenger vehicles to accelerate to 300 miles per hour. They do not
25 need a horizontal agreement among the major car manufacturers to make that choice: engineering
26 constraints, internal safety standards, product-liability exposure, government regulation through
27 the National Highway Traffic Safety Administration, and basic competitive self-interest each
28



1 independently supply the guardrail. The same logic applies here. Each Defendant is free, and in-
2 deed obligated, to set its own safety thresholds, retain its own evaluators, build its own product
3 guardrails, and restrain its own development. None of that requires an agreement with rivals.

4 10. Defendants are also well-positioned to lead on environmental responsibility, and to do
5 so on their own. AI's energy footprint is substantial and growing: training a single large frontier
6 model can consume tens of millions of kilowatt-hours of electricity, and the International Energy
7 Agency projects that, by 2030, global data-center electricity demand will more than double from
8 2024 levels, driven in significant part by AI workloads. Each Defendant has pledged environmen-
9 tal commitments. Each has the engineers, the capital, and the competitive incentive to pursue more
10 efficient training architectures, cleaner energy procurement, and smarter compute scheduling. In-
11 novation in AI efficiency—models that do more with less energy—is itself a domain of competi-
12 tive advantage. An agreement to slow capability development does nothing to make AI more en-
13 vironmentally sound; only engineering competition, applied to efficiency as a prize, will.

14 11. In short, the path to responsible, safe, and environmentally conscious AI development
15 is each company doing the right thing: investing in safety research, adopting rigorous internal
16 evaluations, reducing its environmental footprint, and welcoming government oversight and reg-
17 ulation. That path is fully open to every Defendant. What the antitrust laws forbid is the shortcut
18 of competitors agreeing among themselves to substitute collective restraint for individual account-
19 ability. This action is brought in defense of the competitive process that makes genuine responsi-
20 bility and genuine progress possible.

21 12. The antitrust laws do not permit competitors to decide among themselves that compe-
22 tition is too dangerous. Whether frontier AI should develop more slowly is a question for each
23 company acting alone, or for Congress and the agencies. It is not a question that four competitors
24 may answer together. The law leaves each Defendant free to be as cautious, as responsible, and as
25 environmentally conscious as it chooses to be, just not through agreement with the companies it
26 competes against.



1 17. Plaintiff Cheyenne Hunt is a resident of California. During the Class Period, she per-
2 sonally purchased a paid individual consumer subscription to Claude, ChatGPT, Grok, and Gem-
3 ini. She continues to subscribe and intends to continue subscribing.

4 18. Plaintiff Christine Bullock is a resident of California. During the Class Period, she per-
5 sonally purchased a paid individual consumer subscription to Claude. She continues to subscribe
6 and intends to continue subscribing.

7 19. Plaintiff Nick Spetsas is a resident of Florida. During the Class Period, he personally
8 purchased a paid individual consumer subscription to Claude, ChatGPT, Grok, and Gemini. He
9 continues to subscribe and intends to continue subscribing.

10 20. Plaintiffs purchased their subscriptions to obtain access to the respective Defendant's
11 most capable models and to the features and improvements that Defendant continually releases,
12 and each Plaintiff expected that the Defendant would decide independently, in competition with
13 its rivals, how aggressively to develop and improve its product.

14 21. Plaintiffs have paid, and continue to pay, subscription prices for products whose quality
15 and rate of improvement have been restrained by the agreement alleged herein. Because Plaintiffs
16 continue to subscribe, they face continuing injury from the agreement and have standing to seek
17 injunctive relief.

18 **B. Defendants**

19 22. Defendant Anthropic, PBC ("Anthropic") is a public benefit corporation with its prin-
20 cipal place of business in San Francisco, California. Anthropic develops, markets, and sells Claude
21 and related products and services throughout the United States.

22 23. Anthropic sells consumers paid Claude subscriptions, including Claude Pro and Claude
23 Max, which provide greater usage and access to Anthropic's most advanced models and features.

24 24. Dario Amodei is Anthropic's Chief Executive Officer and a co-founder. Jared Kaplan
25 is a co-founder and its Chief Science Officer. Each spoke and acted with respect to the matters
26 alleged herein as an authorized representative of Anthropic, and each of their statements alleged
27 herein is attributable to Anthropic.
28



1 25. Defendant OpenAI OpCo, LLC (“OpenAI”) is a Delaware limited liability company
2 with its principal place of business in San Francisco, California. It provides ChatGPT to individual
3 consumers under OpenAI’s Terms of Use.

4 26. OpenAI sells consumers paid ChatGPT subscriptions, including ChatGPT Plus (cur-
5 rently approximately \$20 per month) and the higher-priced ChatGPT Pro, which provide greater
6 access to OpenAI’s most advanced models, reasoning, research, coding, image-generation,
7 memory, and other features than OpenAI’s free tier.

8 27. Sam Altman is OpenAI’s Chief Executive Officer. Jakub Pachocki is its Chief Scien-
9 tist. Chris Lehane is its Global Policy Chief. Each spoke and acted with respect to the matters
10 alleged herein as an authorized representative of OpenAI, and each of their statements alleged
11 herein is attributable to OpenAI.

12 28. Defendant SpaceXAI LLC (“SpaceXAI”) is a Nevada limited liability company with a
13 registered office in Austin, Texas. It provides Grok to individual consumers under its Consumer
14 Terms of Service.

15 29. SpaceXAI develops, markets, and sells Grok throughout the United States, including
16 paid SuperGrok subscriptions that increase usage limits and provide access to its frontier models
17 and additional features.

18 30. Elon Musk founded the xAI business, controls it, and speaks for it publicly on the de-
19 velopment and pacing of frontier AI. His statements alleged herein were made as SpaceXAI’s
20 controlling person and authorized representative and are attributable to SpaceXAI.

21 31. Defendant Google LLC (“Google”) is a Delaware limited liability company with its
22 principal place of business in Mountain View, California. Google develops Gemini, including
23 through its Google DeepMind unit, and markets and sells to consumers throughout the United
24 States paid Google AI subscriptions, including Google AI Pro and Google AI Ultra, that provide
25 premium access to Gemini’s most advanced models and capabilities.

26 32. Demis Hassabis is a co-founder and Chair of Google DeepMind and Alphabet’s Chief
27 Scientist. Shane Legg is a co-founder of Google DeepMind and its Chief AGI Scientist. Each
28



1 spoke and acted with respect to the matters alleged herein as an authorized representative of
2 Google, and each of their statements alleged herein is attributable to Google.

3 **C. Agents and Co-Conspirators**

4 33. Whenever this Complaint refers to an act, statement, or agreement of a Defendant, it
5 means that the Defendant acted through its officers, directors, employees, representatives, or
6 agents within the course and scope of their authority.

7 34. Various persons and entities not named as Defendants, including affiliates of Defend-
8 ants and other frontier-AI developers, participated in, facilitated, or knew of aspects of the conduct
9 alleged herein. Plaintiffs reserve the right to amend this Complaint to add them as discovery iden-
10 tifies their roles.

11 35. Each Defendant acted as a co-conspirator with respect to the agreement alleged herein,
12 and each is liable for the acts of the others in furtherance of it.

13 **JURISDICTION, DIVISIONAL ASSIGNMENT, AND VENUE**

14 **A. Jurisdiction**

15 36. This action arises under Section 1 of the Sherman Act, 15 U.S.C. § 1, and Sections 4
16 and 16 of the Clayton Act, 15 U.S.C. §§ 15 and 26.

17 37. This Court has subject-matter jurisdiction under 28 U.S.C. §§ 1331 and 1337(a) and 15
18 U.S.C. §§ 15 and 26.

19 38. This Court also has jurisdiction under the Class Action Fairness Act, 28 U.S.C. §
20 1332(d), because the proposed Class has more than 100 members, at least one Class member is a
21 citizen of a state different from at least one Defendant, and the amount in controversy exceeds
22 \$5,000,000, exclusive of interest and costs.

23 39. Defendants' conduct occurred in and substantially affected interstate commerce. De-
24 fendants sell subscriptions to consumers in every state, receive subscription payments across state
25 lines, deliver their services over the internet from computing and data-center infrastructure located
26 in multiple states, and compete with one another throughout the United States.



B. Divisional Assignment

40. This action is properly assigned to the San Francisco Division under Civil L.R. 3-2(d) because a substantial part of the events or omissions giving rise to the claims occurred in San Francisco County. Anthropic and OpenAI maintain their principal places of business in San Francisco, and the agreement alleged herein was proposed, accepted, and implemented in substantial part by executives and personnel working in and from this District.

C. Personal Jurisdiction and Venue

41. This Court has personal jurisdiction over each Defendant because Anthropic, OpenAI, and Google are at home in this District; each Defendant markets and sells subscriptions to residents of this District and directs the conduct alleged herein at them; and Section 12 of the Clayton Act, 15 U.S.C. § 22, authorizes nationwide service of process on a corporate defendant in an antitrust action brought in a proper district.

42. Venue is proper in this District under 15 U.S.C. § 22 and 28 U.S.C. § 1391(b) and (c) because Anthropic, OpenAI, and Google inhabit and are found in this District; SpaceXAI transacts substantial business in this District, including the sale of subscriptions to its residents; and a substantial part of the events or omissions giving rise to the claims occurred in this District.

FACTUAL ALLEGATIONS**A. Defendants Compete Through the Pace of Innovation**

43. ChatGPT, Claude, Grok, and Gemini are close substitutes for millions of consumers who use them for research, writing, coding, reasoning, mathematical and document analysis, image generation, voice interaction, planning, education, professional work, and automated or agent tasks.

44. Defendants compete principally on capability: models that reason better, code better, answer more accurately, use tools more effectively, process more information, respond faster, perform more complex tasks, and act more autonomously. They also compete on usage limits, context windows, memory, integrations, voice, image and video generation, deep research, agents, connectors, reliability, and speed of access to newly released frontier models.



1 45. Defendants’ paid subscription tiers monetize those differences directly. OpenAI mar-
2 kets paid ChatGPT subscriptions as providing greater access to advanced models, reasoning, re-
3 search, coding, image generation, and memory. Anthropic markets Claude Pro and Max through
4 greater usage, access to advanced models and features, priority availability, and early access to
5 improvements. SpaceXAI markets SuperGrok as providing access to frontier models, greater us-
6 age, and priority access. Google markets Google AI Pro and Ultra as providing access to advanced
7 Gemini models and capabilities unavailable or limited at its free tier.

8 46. The pace of improvement is therefore a central dimension of price and quality compe-
9 tition among Defendants. Each major model release or feature improvement causes consumers to
10 subscribe, cancel, switch products, maintain multiple subscriptions, or change which product they
11 use most. Defendants closely monitor one another’s releases, benchmark results, features, usage
12 limits, and pricing, and each responds when a rival ships a materially better model.

13 47. Before the challenged agreement, that rivalry drove each Defendant to invest aggres-
14 sively in research, compute, training runs, engineering, and deployment. Defendants themselves
15 have repeatedly described the sector as a race and acknowledged that the commercial pressure it
16 generates makes unilateral slowing costly.

17 48. The stakes compound. OpenAI researchers have explained that a principal research
18 priority is using advanced AI to accelerate AI research itself, and Anthropic describes recursive
19 self-improvement as a principal driver of accelerating capability. An agreement about today’s pace
20 of development therefore shapes the trajectory and timing of every subsequent generation of prod-
21 ucts.

22 **B. Defendants Recognized That No Firm Would Slow Alone**

23 49. In July 2026, senior executives of Anthropic, OpenAI, and Google published a state-
24 ment titled Pacing the Frontier. Signatories included Amodei; Kaplan; Pachocki; OpenAI Chief
25 Research Officer Mark Chen; Legg; and numerous other employees of Anthropic, OpenAI, and
26 Google.



1 50. The statement identified the collective-action problem at the heart of this case: each
2 company was under “intense competitive pressure not to unilaterally slow” capability develop-
3 ment. It called for tools capable of deliberately pacing frontier-wide development.

4 51. Altman told reporters in July that OpenAI had helped shape the statement’s language
5 and that he had discussed the “need” to pace AI development with White House officials as models
6 became more capable.

7 52. The statement thus reflected a shared understanding among Defendants’ leadership,
8 months before September 12, that slowing was commercially costly for any firm that acted alone,
9 and that the cost could be avoided only if the frontier laboratories slowed together.

10 **C. Defendants Built the Machinery for Coordination Before September 12**

11 53. On July 14, 2026, Hassabis proposed a U.S.-led standards body for frontier AI modeled
12 in part on the Financial Industry Regulatory Authority. Under the proposal, frontier laboratories
13 would initially submit models voluntarily for pre-release review; the body could later impose com-
14 mon requirements before models entered the U.S. market and could escalate those requirements
15 as risks increased. Contemporaneous reporting described the proposal as a mechanism that could
16 coordinate an industry-wide slowdown.

17 54. Such a body would give competing developers a standing forum in which to exchange
18 information, adopt common requirements, review one another’s frontier systems, and coordinate
19 responses to capability thresholds.

20 55. Beginning in July 2026, representatives of Anthropic, OpenAI, and Google below the
21 CEO level formed a working group that met regularly to work together on creating a standards
22 body for the industry. On September 13, The Information reported that the working group’s meet-
23 ings were ongoing and that it had met during the preceding week. On September 15, OpenAI
24 confirmed that it had been discussing AI safety with Anthropic and Google DeepMind for several
25 weeks.

26 56. During the same period, Altman told an OpenAI companywide town hall that the major
27 laboratories might need to create a standards body themselves if government did not, and that the
28



1 best course could be to act together with the other laboratories. According to published reports,
2 Altman also told OpenAI employees that OpenAI might slow the pace of frontier development,
3 potentially in conjunction with other laboratories.

4 57. On September 6, 2026, OpenAI published on its website an essay by Pachocki titled
5 An Alien Mind. Pachocki wrote that one of the principal options available to frontier developers
6 was “coordinating to slow down future development”; that no laboratory had solved alignment and
7 monitoring well enough to continue scaling at maximum speed indefinitely; and that he expected
8 voluntary slowdowns to become common until shared safety thresholds were established.

9 58. On September 10, 2026, WIRED reported that OpenAI had sought guidance from
10 members of Congress on whether coordinating an industry-wide slowdown in frontier-AI devel-
11 opment could violate the antitrust laws. OpenAI did not substantively deny the report. The inquiry
12 reflected OpenAI’s understanding that the contemplated conduct was coordination among com-
13 petitors, not unilateral safety decision-making.

14 59. On September 11, 2026, in an interview with Fortune, Altman was asked why leading
15 AI executives, including Amodei and Musk, did not get together and develop a common plan. He
16 answered: “I think that will happen.” He declined to “pre-announce private discussions” that he
17 believed should later be shared “as a group.”

18 60. The day before Amodei’s essay, Altman thus referred to private discussions about in-
19 dustry collaboration, declined to disclose their details, and said he thought a common plan would
20 happen.

21 **D. Amodei Proposes the Agreement: “We Must Pace the Frontier”**

22 61. On the morning of September 12, 2026, Amodei published We Must Pace the Frontier
23 and promoted it publicly.

24 62. The essay’s thesis was direct: “We must slow the pace at which we improve the capa-
25 bilities of AI models.”

26 63. Amodei proposed three steps. The first—embedding third-party evaluators with em-
27 ployee-like access to models and training pipelines—Anthropic adopted unilaterally.
28



1 64. The second step, which Amodei called “Democratic Coordination,” he expressly dis-
2 tinguished from the first: it “requires industry-wide coordination.” Under it, the frontier laborato-
3 ries would jointly establish common safety standards and “limits on the rate of unchecked AI pro-
4 gress.”

5 65. Amodei identified the inputs through which those limits would operate. He proposed
6 consideration of limits on training compute, on the nature of training runs, and on the internal use
7 of AI to improve AI—the principal means by which each Defendant improves its products. He
8 proposed capability “checkpoints” at which continued advancement would depend on agreed
9 safety certifications or conditions.

10 66. Amodei also proposed the means of enforcement. Embedded evaluators, he explained,
11 would make “verifiable pacing” practical by giving outsiders access to detailed properties of mod-
12 els and training pipelines—that is, the ability to confirm that each participant was in fact slowing
13 as agreed.

14 67. Amodei was explicit about the commercial problem the plan solved. A company that
15 slows alone while rivals race ahead loses customers, revenue, talent, and technological leadership;
16 coordinated pacing would let frontier developers slow “without sacrificing commercial ad-
17 vantage.” That is the economic function of every output-restricting cartel: to remove the competi-
18 tive penalty for restraint by ensuring that rivals restrain themselves too.

19 68. Amodei acknowledged that legislation might take too long and urged AI companies, in
20 parallel with any government action, to “voluntarily work together” on the proposed arrangements.

21 69. Amodei also acknowledged the antitrust problem. His essay stated that government
22 mediation would be helpful and that the government would need to issue a narrow waiver for
23 certain kinds of safety conversations. No such waiver exists. Congress has enacted no exemption,
24 and no agency has compelled the conduct alleged herein.

25 70. The proposal thus supplied every element of a horizontal agreement in restraint of
26 trade: the parties (the frontier laboratories); the term (a slower rate of capability advancement than
27 each would choose alone); the mechanisms (shared limits on compute, training runs, and self-
28



1 improvement, and capability checkpoints); and the means of policing defection (verifiable pacing
2 through embedded evaluators).

3 **E. Defendants Immediately Accept**

4 71. Within approximately an hour of publication, Musk quote-posted Amodei's announce-
5 ment of the proposal and responded: "Dario is right."

6 72. Later that day, Altman wrote: "I agree with Dario that we need to pace the frontier."
7 He added that pacing had been a primary subject of discussion inside OpenAI in recent weeks;
8 committed OpenAI to Amodei's first step, stating that OpenAI "will do the same" in giving inde-
9 pendent evaluators employee-like access; and said OpenAI would have more to share soon.

10 73. Also on September 12, Hassabis stated publicly that Amodei's essay pointed toward
11 "the right path forward" and that "the direction is correct," while noting that implementation details
12 required further work. He tied his endorsement to his own proposed "industry-wide standards body
13 for frontier AI"—the mechanism the Anthropic-OpenAI-Google working group had been devel-
14 oping since July.

15 74. By the end of September 12, therefore, a senior executive or controlling person of each
16 Defendant had publicly signaled assent to a proposal that the frontier laboratories coordinate to
17 slow the pace of capability development. Each signal was given with knowledge of the others, in
18 sequence, on a single day, in direct response to the proposal.

19 75. A public offer followed by public acceptances forms an agreement no less than more
20 traditional cartel behavior of private, secretive communications. Conducting part of the exchange
21 in view of the market assured each participant that its rivals were committed.

22 **F. Defendants Confirm the Agreement and Its Object**

23 76. On September 14, Altman published further statements explaining what OpenAI had
24 agreed to. OpenAI favored consistent frontier-AI safety rules and common approaches to safety
25 and monitoring across the industry, and looked forward to "collaborating with our colleagues
26 across the industry" to develop them.



1 77. Altman then stated the object of the agreement in plain terms: AI progress “should be
2 slower than it otherwise could be.”

3 78. Altman added that OpenAI did not intend to wait for an antitrust exemption or legisla-
4 tion before beginning this work, and that the firms should first do what they could themselves.

5 79. On September 15, OpenAI Global Policy Chief Chris Lehane confirmed that OpenAI
6 had been working with Anthropic and Google DeepMind on these issues for several weeks—be-
7 fore Amodei’s essay—and that OpenAI believed the firms could work together without waiting
8 for an antitrust waiver.

9 80. Lehane’s statement corroborated the reporting that OpenAI, Anthropic, and Google
10 had been communicating privately before the public exchange of September 12, and established
11 that the September 12 statements were the public face of an existing course of dealing among
12 competitors.

13 81. The coordination continues. Amodei continues to advocate interfirm discussions on the
14 remaining stages of his framework; Altman has said that OpenAI anticipates further collaboration
15 and additional announcements; and the working group and the proposed standards body remain
16 the vehicles through which Defendants intend to implement pacing.

17 **G. Each Defendant Knowingly Joined the Agreement**

18 82. Anthropic proposed the agreement through its Chief Executive Officer, defined its
19 terms, supplied its enforcement mechanism, participated in the working group from July, and has
20 continued to advocate implementation.

21 83. OpenAI participated in the working group from July; helped draft the July statement;
22 explored through its CEO—with employees and with the White House—whether to coordinate a
23 slowdown with other laboratories, and sought guidance from members of Congress about such
24 coordination; published its Chief Scientist’s endorsement of “coordinating to slow down future
25 development” six days before the proposal; accepted the proposal through its CEO on September
26 12; stated its object on September 14; and confirmed weeks of prior collaboration through its
27 Global Policy Chief on September 15.



1 84. Google participated in the working group from July; proposed, through Hassabis, the
2 standards body that Defendants have identified as the forum for coordination; had a Google Deep-
3 Mind co-founder sign the July statement; and, on September 12, endorsed the direction of Amo-
4 dei's proposal while directing its implementation toward the body Google had proposed.

5 85. SpaceXAI, through the person who founded and controls its AI business, publicly en-
6 dored the proposal within an hour of its publication, communicating to its three principal rivals
7 that SpaceXAI would not exploit their restraint by racing ahead. Altman's September 11 acknowl-
8 edgment of private discussions came in direct response to a question that named Musk as a partic-
9 ipant in the contemplated common plan.

10 **H. The Agreement Is Not Mere Parallel Conduct**

11 86. The agreement is evidenced by direct communications among Defendants: a written
12 proposal for "industry-wide coordination" to limit the rate of AI progress and innovation; express
13 statements of agreement by the proposer's competitors; and admissions that those competitors had
14 been working together privately for weeks.

15 87. The circumstances also include every traditional indicator that Defendants' conduct
16 was the product of agreement rather than independent decisionmaking:

17 (a) Motive. Each Defendant had a strong motive to conspire, because, as Defendants them-
18 selves stated, slowing the pace of innovation was seen to be commercially untenable for any firm
19 acting alone but advantageous if all slowed together.

20 (b) Action against unilateral self-interest. According to Defendants' own statements, absent
21 an agreement, no Defendant would announce that it will improve its products more slowly than it
22 could, because the announcement would invite rivals to take its customers. Each Defendant's an-
23 nouncement is explicable only by its understanding that its rivals had made, or would make, the
24 same commitment.

25 (c) Interfirm communications and opportunity. Representatives of Anthropic, OpenAI, and
26 Google met regularly from July through September; Altman acknowledged private discussions on
27 September 11; and Lehane confirmed weeks of collaboration on September 15.



(d) Market structure. Frontier AI is highly concentrated among a handful of firms protected by enormous entry barriers, and Defendants have repeatedly described themselves as the “frontier labs” whose joint participation is necessary for pacing to work.

(e) Consciousness of antitrust risk. OpenAI asked members of Congress whether coordinating a slowdown would violate the antitrust laws; Amodei wrote that the government would need to issue a narrow waiver for certain kinds of safety conversations; Altman announced that OpenAI would proceed without one.

(f) Pretext. Every safety objective Defendants have identified can be pursued unilaterally, through independent evaluators, or through government regulation. Only the elimination of competitive pressure requires an agreement among rivals.

I. Plaintiffs Do Not Challenge Defendants’ Petitioning of Government

88. Plaintiffs do not seek to impose liability for any Defendant’s advocacy to Congress, the White House, or any agency, including advocacy for AI regulation, for a standards body created by statute, or for an antitrust exemption.

89. The restraint challenged here is Defendants’ private agreement among themselves, formed and implemented outside any governmental process, to slow the pace of competitive development. Defendants’ own statements establish that they are proceeding with that agreement without waiting for, and independent of, any governmental action.

90. Defendants’ petitioning statements are alleged only as evidence of Defendants’ knowledge, intent, and understanding of the competitive problem the agreement addresses, and of the antitrust risk they knowingly accepted.

J. How the Agreement Restrains Competition

91. In a competitive market, each Defendant independently decides how much to invest, how quickly to train new models, when to begin training runs, how much compute to deploy, when to release improvements, and how aggressively to pursue new capabilities.

92. Under the agreement, each Defendant makes those decisions knowing that its principal rivals have agreed to reduce their own pace. That knowledge eliminates the competitive penalty



1 for slowing—the risk that a rival will race ahead—and allows each Defendant to reduce its invest-
2 ment and competitive effort without losing ground.

3 93. A cartel can restrain competition by reducing output, quality, investment, or competi-
4 tive effort; it need not eliminate them. Here the agreement operates through a common commit-
5 ment to reduce the rate at which competing products improve, shared checkpoints and standards
6 that delay advancement, reciprocal expectations about development conduct, and monitoring that
7 deters defection.

8 94. Plaintiffs challenge Defendants’ agreement to coordinate those decisions. Each De-
9 fendant remains free to make them unilaterally.

10 MARKET POWER

11 95. Defendants’ agreement is a naked horizontal restraint on a central dimension of com-
12 petition among direct rivals, and it is unlawful without proof of a relevant market or market power.

13 96. To the extent market definition is required, the relevant product market is paid con-
14 sumer subscriptions to general-purpose frontier generative-AI assistants (the “Paid Frontier AI
15 Assistant Subscription Market”). Products in this market provide consumers ongoing access,
16 through a conversational or assistant interface, to highly capable, general-purpose generative-AI
17 systems that perform a wide range of intellectual and creative tasks. The principal products are the
18 paid tiers of ChatGPT, Claude, Grok, and Gemini.

19 97. These products are reasonably interchangeable. Consumers use them for the same
20 tasks, compare them on the same attributes—model capability, reasoning, coding, accuracy, re-
21 search functionality, speed, multimodality, agents, context, integrations, usage limits, and price—
22 and switch among them in response to differences in capability and price. Defendants market them
23 as alternatives to one another.

24 98. Conventional search engines, ordinary productivity software, narrow machine-learning
25 tools, and human professional services are not reasonable substitutes. None provides general-pur-
26 pose frontier capability through an assistant interface, and consumers would not substitute to them



1 in response to a small but significant, non-transitory decrease in the quality or rate of improvement,
2 or increase in the price, of paid frontier-AI subscriptions.

3 99. Free tiers of Defendants' own products do not constrain the paid market. They impose
4 substantially lower usage limits, restrict or exclude access to frontier models, and offer fewer fea-
5 tures and lower priority. And because Defendants control their free tiers as well, the agreement
6 restrains the improvement of free and paid tiers alike; a free tier offers no competitive discipline
7 that its owner has agreed to withhold.

8 100. In the alternative, the agreement restrains competition in an innovation market for the
9 development of general-purpose frontier generative-AI models, in which Defendants are the prin-
10 cipal participants and on which the restraint operates directly.

11 101. The relevant geographic market is the United States. Defendants sell subscriptions
12 nationwide on uniform terms, and consumers cannot practically obtain frontier-AI subscriptions
13 from sellers outside the country on materially different terms.

14 102. The Paid Frontier AI Assistant Subscription Market is highly concentrated. On infor-
15 mation and belief, Defendants collectively account for at least 80% of paid consumer subscriptions
16 to general-purpose frontier generative-AI assistants in the United States.

17 103. Entry at the frontier requires expenditures on computing infrastructure, specialized
18 semiconductors, electricity, data-center capacity, engineering, research personnel, model training,
19 safety systems, and deployment infrastructure; access to scarce technical talent and large-scale
20 compute; consumer distribution and product reliability; and the ability to fund successive genera-
21 tions of models. These barriers limit the number of firms capable of imposing competitive disci-
22 pline at the frontier to a handful.

23 104. Defendants' own statements confirm that they regard one another as the relevant com-
24 petitive set. Their public discussions refer to a handful of "frontier labs," to coordination among
25 those laboratories, and to the inability of any one of them to slow without competitive conse-
26 quence. In a market this concentrated, an agreement among Defendants determines market-wide
27 competitive incentives.



ANTICOMPETITIVE EFFECTS

105. The agreement has the purpose, tendency, and effect of restraining competition in the development and improvement of frontier generative-AI products.

106. The agreement reduces the independent competitive pressure that otherwise drives each Defendant to improve its products as rapidly as it can, and it reduces each Defendant's incentive to devote research effort and compute to capability advancement, to begin or expand training runs, to accelerate model development, to deploy innovations promptly, and to respond aggressively to rivals' advances.

107. The agreement causes and threatens delay in the development and delivery of improved capabilities to consumers; reduced product quality relative to what unrestrained competition would produce; reduced investment and competitive effort; and fewer or later improvements in reasoning, coding, research, multimodal capability, reliability, agents, speed, context, memory, and other features.

108. The agreement diminishes the value consumers receive for the subscription prices they pay. Consumers pay the same or higher nominal prices for products that improve more slowly than they would under unrestrained competition; the quality-adjusted price of Defendants' subscriptions is therefore higher than it would be absent the agreement.

109. The agreement reduces consumer choice by foreclosing the possibility that one competitor will break from the pack and pursue technological improvement more aggressively than its rivals, and it harms the competitive process itself by substituting the collective judgment of competitors for independent rivalry.

110. Because the agreement was formed recently and Defendants' development cycles run months, its full effect on released products has not yet manifested. But the agreement has already altered each Defendant's incentives and, on information and belief, its investment, training, and release decisions, and its continuing operation threatens exactly the injury that Section 16 of the Clayton Act exists to prevent.



111. The agreement has no cognizable procompetitive justification. It is not reasonably necessary to permit any Defendant to conduct safety research, test its own models, adopt its own safety standards, retain independent evaluators, disclose safety incidents, or slow its own development. Each Defendant can do all of those things unilaterally, without limitation.

112. To the extent Defendants contend that competition in frontier AI is itself unsafe and should be curtailed, that contention is not a procompetitive justification cognizable under the Sherman Act. The Act does not permit competitors to decide among themselves that competition is undesirable as a policy matter; that judgment belongs to Congress.

113. Nor is the agreement ancillary to any legitimate collaboration among Defendants. They do not integrate productive assets, share risk, or jointly produce anything. They remain separate firms that independently develop, market, price, and sell competing products, and the sole function of their agreement is to reduce the competition among them.

114. No statute exempts Defendants or their agreement to “pace the frontier” from the antitrust laws, and no governmental body has compelled it.

ANTITRUST INJURY

115. Plaintiffs and Class members purchase directly from Defendants the products whose improvement the agreement restrains. They pay Defendants money for access to products whose value depends on technological capability and continued improvement.

116. Absent the agreement, each Defendant would face stronger independent incentives to invest, train, improve, and release capabilities rapidly, because each would face the risk that a rival would do so first, and vigorous competition would force Defendants to provide more value for the subscription prices they charge.

117. As a direct and foreseeable result of the agreement, Plaintiffs and Class members have paid and continue to pay subscription prices that exceed the competitive value of the products they receive. Stated another way, they receive less capability, slower improvement, or both, for the prices they pay. That difference is an overcharge recoverable under Section 4 of the Clayton Act.



118. That injury is measurable on a class-wide basis through common evidence, including Defendants' development and release plans; actual and counterfactual investment, compute, and training decisions; product capabilities and benchmark performance; subscription prices; consumer demand, switching, and usage data; and accepted economic methods for valuing model and feature improvements. The amount of damages will be established through discovery and expert proof.

119. Plaintiffs' injuries are injuries of the type the antitrust laws were designed to prevent, and they flow directly from what makes Defendants' conduct unlawful: horizontal competitors agreed to diminish the competitive pressure to improve the products Plaintiffs buy.

120. Plaintiffs and Class members are direct purchasers and the most efficient enforcers of the antitrust laws with respect to this agreement. No more direct group of purchasers exists.

121. Because Plaintiffs and Class members continue to purchase Defendants' subscriptions, and because the agreement is ongoing, they face threatened loss and damage from its continuation and are entitled to injunctive relief under Section 16 of the Clayton Act.

122. Defendants are jointly and severally liable for all damages caused by the agreement.

CLASS ACTION ALLEGATIONS

123. Plaintiffs bring this action individually and on behalf of the following class under Federal Rules of Civil Procedure 23(a), 23(b)(2), and 23(b)(3):

Nationwide Class: All persons in the United States who, from September 12, 2026 through the date on which the effects of Defendants' unlawful conduct cease (the "Class Period"), purchased directly from one or more Defendants a paid individual consumer subscription providing premium access to ChatGPT, Claude, Grok, or Gemini.

124. Plaintiffs also propose, as appropriate, the following Defendant-specific Subclasses:

OpenAI Subscriber Subclass: All persons in the United States who, during the Class Period, purchased directly from OpenAI a paid individual consumer subscription to ChatGPT.

Anthropic Subscriber Subclass: All persons in the United States who, during the Class Period, purchased directly from Anthropic a paid individual consumer subscription to Claude.



1 **SpaceXAI Subscriber Subclass:** All persons in the United States who, during the Class
2 Period, purchased directly from SpaceXAI a paid individual consumer subscription to Grok.

3 **Google Subscriber Subclass:** All persons in the United States who, during the Class Pe-
4 riod, purchased directly from Google a paid individual consumer subscription providing premium
5 access to Gemini.

6 125. Plaintiffs reserve the right to modify the Class and Subclass definitions based on facts
7 developed in discovery, including evidence that the agreement was formed before September 12,
8 2026.

9 126. Excluded from the Class are Defendants; their parents, subsidiaries, and affiliates;
10 their officers and directors; governmental entities; the Court and its staff assigned to this action;
11 and any person or entity that validly and timely requests exclusion from a Rule 23(b)(3) Class.

12 **Numerosity**

13 127. The members of the Class and each Subclass are so numerous that joinder is imprac-
14 ticable. Each Defendant sells paid subscriptions to large numbers of consumers throughout the
15 United States, and Class members can be identified from Defendants' subscription and billing
16 records using objective criteria.

17 **Commonality**

18 128. This action presents questions of law and fact common to the Class, including:

19 (a) whether Defendants entered a contract, combination, or conspiracy concerning the pace
20 of frontier-AI development;

21 (b) the scope, participants, terms, and duration of that agreement;

22 (c) whether Defendants communicated privately before or after their September 12, 2026
23 public statements, and what they agreed to;

24 (d) whether the agreement is a naked horizontal restraint that is unlawful per se, unlawful
25 under abbreviated analysis, or unlawful under the rule of reason;

26 (e) whether a relevant market must be defined and, if so, its boundaries, and whether De-
27 fendants possess market power individually or collectively;



1 (f) whether the agreement reduced competition in innovation, quality, output, investment,
2 or product development;

3 (g) whether the agreement caused paid AI subscriptions to provide less value than they
4 would have provided absent the restraint;

5 (h) whether antitrust impact and damages can be measured using common economic evi-
6 dence;

7 (i) whether Defendants' conduct is ongoing or threatens continuing injury; and

8 (j) the nature and scope of appropriate injunctive relief.

9 129. These questions will be resolved through evidence common to the Class, including
10 Defendants' communications, meeting records, internal development plans, research budgets,
11 compute allocations, model-training plans, release schedules, competitive analyses, subscription
12 and pricing records, and common economic evidence.

13 **Typicality**

14 130. Plaintiffs' claims are typical of the claims of the Class and any applicable Subclass.
15 Plaintiffs purchased a subscription product affected by the same agreement, suffered the same type
16 of antitrust injury, and will confront the same defenses concerning the existence, legality, and
17 effect of the agreement.

18 **Adequacy**

19 131. Plaintiffs will fairly and adequately protect the interests of the Class. Plaintiffs' inter-
20 ests do not conflict with those of absent Class members, and Plaintiffs have retained counsel ex-
21 periented in complex antitrust class actions who are capable of prosecuting this action vigorously.

22 **Rule 23(b)(3)**

23 132. Common issues predominate over individual issues. The existence and scope of the
24 agreement, Defendants' common course of conduct, the nature of the restraint, and its competitive
25 effects will be established principally through evidence common to all Class members, and anti-
26 trust impact and damages are susceptible to proof through methodologies applicable to the Class
27 or Subclasses as a whole.
28



1 133. A class action is superior to individual litigation. Individual consumers' damages are
2 too small to justify separate antitrust suits requiring extensive discovery of some of the world's
3 largest and most sophisticated technology companies, and concentrating the litigation in one pro-
4 ceeding promotes efficiency, avoids inconsistent adjudications, and serves the private-enforce-
5 ment purposes of the antitrust laws.

6 **Rule 23(b)(2)**

7 134. Defendants have acted and continue to act on grounds generally applicable to the
8 Class, and their ongoing coordination threatens continuing injury to every consumer who pur-
9 chases their subscriptions. Final injunctive and declaratory relief concerning the agreement is
10 therefore appropriate respecting the Class as a whole.

11 **CLAIMS FOR RELIEF**

12 **FIRST CLAIM FOR RELIEF**

13 **Violation of Section 1 of the Sherman Act, 15 U.S.C. § 1**

14 **(Against All Defendants)**

15 135. Plaintiffs incorporate by reference the preceding paragraphs.

16 136. Defendants are horizontal competitors in the development and sale of general-purpose
17 frontier generative-AI products and of paid consumer subscriptions to those products, and each is
18 a separate economic actor capable of conspiring with the others.

19 137. Beginning no later than September 12, 2026, and on information and belief earlier,
20 Defendants entered into a contract, combination, or conspiracy to restrain the pace at which they
21 develop and improve their competing frontier-AI products.

22 138. Defendants formed and manifested the agreement through, among other things, Amo-
23 dei's written proposal for industry-wide coordination to limit the rate of AI progress; Musk's,
24 Altman's, and Hassabis's public statements of assent on September 12, 2026; the private discus-
25 sions and recurring working-group meetings among Anthropic, OpenAI, and Google from July
26 2026 onward; Altman's September 14 statement that AI progress "should be slower than it other-
27 wise could be"; and Lehane's September 15 confirmation of weeks of collaboration.
28



1 139. The agreement restrains competition in model development, product quality, innova-
2 tion, investment, output, and the timing of new capabilities, and its purpose and effect is to replace
3 each Defendant's independent decisions about the pace of innovation with collective restraint.

4 140. The agreement is a naked horizontal restraint on output and product quality among
5 competitors. It is not ancillary to any integration of productive assets or legitimate joint venture;
6 Defendants remain separate firms that independently develop, market, price, and sell competing
7 products. The agreement is therefore unlawful per se.

8 141. Alternatively, the agreement is unlawful under abbreviated, or "quick-look," analysis,
9 because an observer with even a rudimentary understanding of economics would conclude that an
10 agreement among the principal competitors in a market that technological progress in their com-
11 peting products should be slower than competition would otherwise produce has an anticompeti-
12 tive effect on consumers.

13 142. Alternatively, the agreement is unlawful under the rule of reason. Defendants possess
14 substantial collective power in the relevant market alleged above; the agreement has caused and
15 threatens substantial anticompetitive effects, including slower innovation, reduced quality, dimin-
16 ished investment and competitive effort, delayed improvements, and higher quality-adjusted
17 prices; the agreement has no cognizable procompetitive justification; and any legitimate objective
18 Defendants may identify can be achieved through substantially less restrictive means, including
19 unilateral safety decisions, independent third-party evaluation, disclosure, government regulation,
20 safety research, and technical standards that do not limit independent competitive development.

21 143. Defendants' conduct occurred in and substantially affected interstate commerce.

22 144. As a direct and proximate result of Defendants' violation, Plaintiffs and Class mem-
23 bers have been injured in their business or property, in that they have paid subscription prices
24 exceeding the competitive value of the products they received and have received lower-quality or
25 more slowly improving products than competition would have produced.

26 145. Plaintiffs and Class members are entitled to recover treble damages, the costs of suit,
27 and reasonable attorneys' fees under Section 4 of the Clayton Act, 15 U.S.C. § 15.
28



SECOND CLAIM FOR RELIEF

Injunctive Relief Under Section 16 of the Clayton Act, 15 U.S.C. § 26

(Against All Defendants)

146. Plaintiffs incorporate by reference the preceding paragraphs.

147. Defendants' agreement violates Section 1 of the Sherman Act for the reasons alleged above.

148. The agreement is ongoing. Defendants continue to communicate and collaborate concerning frontier-AI standards, safety thresholds, evaluators, and pacing, and they have announced that further collaboration and announcements will follow.

149. Plaintiffs and Class members who continue to purchase Defendants' subscriptions face threatened loss and damage from the continued suppression of independent competition in innovation, and monetary damages alone will not prevent that injury.

150. Plaintiffs therefore seek injunctive relief prohibiting each Defendant from entering into, maintaining, enforcing, or monitoring any agreement with a horizontal competitor concerning:

(a) the rate at which competing AI products are developed, improved, trained, or released;

(b) limits on training compute, or on the nature, frequency, or timing of training runs, adopted pursuant to a horizontal pacing arrangement;

(c) limits on the use of AI to develop improved AI systems, where imposed by agreement among competitors;

(d) coordinated delays in the development or release of models or features;

(e) capability checkpoints that operate as horizontally agreed restraints on competitive development, as distinct from lawful governmental requirements or independent safety measures; and

(f) the exchange of competitively sensitive information for the purpose of policing or enforcing any such restraint.



1 151. Plaintiffs do not seek to prohibit any Defendant from independently adopting safety
2 measures or product guardrails, independently considering and reducing environmental impacts,
3 independently slowing or suspending its own development, independently retaining evaluators,
4 conducting lawful safety research, complying with governmental requirements, petitioning any
5 governmental body, or participating in legitimate standard-setting activities that do not restrain
6 independent competition over the pace of development.

7 152. Plaintiffs and Class members are entitled to injunctive relief under 15 U.S.C. § 26.

8 **PRAYER FOR RELIEF**

9 WHEREFORE, Plaintiffs, individually and on behalf of the proposed Class and Sub-
10 classes, respectfully request that the Court:

11 A. certify this action as a class action under Federal Rule of Civil Procedure 23, appoint
12 Plaintiffs as representatives of the Class and any applicable Subclass, and appoint Plaintiffs' coun-
13 sel as Class Counsel;

14 B. adjudge and decree that Defendants entered into and participated in an unlawful con-
15 tract, combination, or conspiracy in violation of Section 1 of the Sherman Act, 15 U.S.C. § 1;

16 C. award Plaintiffs and the Class damages in an amount to be established at trial, trebled
17 as provided by Section 4 of the Clayton Act, 15 U.S.C. § 15;

18 D. enter preliminary and permanent injunctive relief under Section 16 of the Clayton Act,
19 15 U.S.C. § 26, as set forth above, preventing Defendants from continuing, renewing, enforcing,
20 or implementing any horizontal agreement to restrain competitive innovation;

21 E. award Plaintiffs and the Class pre-judgment and post-judgment interest to the fullest
22 extent permitted by law;

23 F. award Plaintiffs reasonable attorneys' fees and costs of suit as provided by law; and

24 G. grant such other and further relief as the Court deems just and proper.

25 **DEMAND FOR JURY TRIAL**

26 Pursuant to Federal Rule of Civil Procedure 38(b), Plaintiffs demand a trial by jury on all
27 claims and issues so triable.
28



1 Dated: September 18, 2026

Respectfully submitted,

2 /s/ Andrew T. Tutt
3 Nicholas C. Rowley (SBN 220036)
4 Lead Counsel
5 **TRIAL LAWYERS FOR JUSTICE**
6 8235 Santa Monica Blvd.
7 West Hollywood, CA 90046
8 Tel.: (310) 855-3727
9 Fax: (310) 855-3595
10 ncr@tl4j.com

11 Andrew T. Tutt (SBN 343206)
12 R. Stanton Jones (DC Bar No. 987088)*
13 **TRIAL LAWYERS FOR JUSTICE**
14 421 W. Water Street
15 Decorah, IA 52101
16 Tel.: (202) 490-7080
17 Fax: (563) 204-4205 (Tutt)
18 andrew@tl4j.com
19 stanton@tl4j.com

20 Jakob Z. Norman (WY Bar No. 6-3685)*
21 **TRIAL LAWYERS FOR JUSTICE**
22 724 S. Durbin Street
23 Casper, WY 82601
24 Tel.: (307) 288-4400 ext. 101
25 jakob@tl4j.com

26 *Attorneys for Plaintiffs and the Proposed Class*

27 * Application for *Pro Hac Vice* Admission Forthcoming
28

